

Bayesian Modeling of Polar Sea Ice from IceSat-2 Freeboard Measurements

Jeff Liu

August 30, 2020

1 Introduction

Polar sea ice thickness is of interest to researchers for a variety of reasons including maritime navigation and climate change. In 2018, NASA launched the IceSat-2 satellite, which relies a laser altimeter to help study polar sea ice. To measure the thickness of this ice, the satellite takes measurements of “freeboard”, which is the height of the ice above sea water. Typically, $7/8$ of the ice lies below water, so the freeboard can provide direct inference for the true thickness of the ice; however, this measurement is complicated by an unknown amount of snow. Any snow present may weigh down the ice, pushing it deeper into the water, which renders this $7/8$ rule of thumb inaccurate.

In response to this complication, scientists have devised a more sophisticated model that incorporates the amount of snow present into the relationship between ice thickness and freeboard. It relies on three additional parameters: the densities of the surrounding sea water, the sea ice in question, and the snow on top of the ice. My goal in this paper is to incorporate this physical model as well as prior knowledge of the parameters to estimate the sea ice thickness corresponding to a freeboard measurement of 0.1 meters. I will proceed to evaluate this estimate using Bayesian inference techniques by defining a prior and likelihood and measuring the posterior.

I will first proceed to elaborate on this physical model and formulate the Bayesian process to evaluate the expected ice thickness in section II, including discussion of the assumptions, prior, hyper priors, forward model, likelihood functions, and posterior densities. Following that, I will discuss the computational strategies for computing the posterior densities, including Markov Chain Monte Carlo, in section III. Finally, I will report my model results and analyze model accuracy in section IV, and analyze and discuss the results in section V.

2 Formulation

2.1 Physical Model

The physical model relating the ice thickness H , snow thickness S , and freeboard H is the following:

$$F = \frac{\rho_w - \rho_i}{\rho_w} \left(H - S \frac{\rho_s}{\rho_w - \rho_i} \right),$$

where ρ_w , ρ_i , and ρ_s are the densities of the sea water, sea ice, and snow, respectively. In the literature, it is given that Arctic seawater density ρ_w ranges from 1010 to 1030 kg/m^3 , sea ice density ρ_i ranges from 750 to 940 kg/m^3 while having an average value of 910 kg/m^3 , and snow density averages 300 kg/m^3 while ranging from 100 to 400 kg/m^3 . In addition, I add a constraint on sea ice thickness from Dr. Matthew Perno's Math 76 lecture on a similar model, where sea ice thickness ranges from 0.1 to 5 meters while having an average value of 2 meters. No information is given on snow thickness, but it is safe to say that snow thickness must be greater than or equal to zero meters.

2.2 Assumptions

To make this inference possible, I will have to make many assumptions because the given information is quite sparse.

First, I must give some characterization to a prior density over these parameters. It is helpful that I have a given range to work with for the parameters, but I will have to arbitrarily assume variances for all of them. The potential bias introduced by of any poorly assumed variances is dampened by the fact that I have a predefined range of possibilities.

In concert with these variance assumptions, I also assume that the prior densities are truncated Gaussian densities. The reason for this choice is twofold. First, Gaussian densities are easily computed, and an added truncation is trivial as long as we do not need to integrate to find the normalizing constant. Second, these truncations allow me to impose a strict range upon the parameters, which helps with limiting the bias of arbitrarily chosen variances in my first assumption.

My third assumption is that this model has additive noise. It is given that IceSat-2 can measure ice height within a few centimeters. Since this measurement affects only the freeboard, which is not multiplied with anything in the physical model, I am assuming that this error is additive. However, I will not make a hyperprior estimate on the exact amount of additive noise, and will reserve that to a hyperprior distribution which I will discuss later.

I will also need to make certain assumptions about the snow depth to make this model more easily computable. I visited a NASA website discussing arctic snow depth over sea ice (<https://earthobservatory.nasa.gov/images/146758/mapping-snow-on-arctic-sea-ice>), and used the graphics to semi-arbitrarily decide an average snow depth of 0.15 meters. However, giving an arbitrary estimate of the snow depth variance may introduce too much bias into the

model since I have no little information. Therefore, I will be modeling this variance as a hyperparameter with a hyperprior as well.

Because I am assuming truncated Gaussians for all prior densities, I will assume that the inverse-Gamma distribution will serve as a proper hyperprior on my additive noise/error variance and my snow thickness variance. This is typically chosen for conjugacy reasons, which I cannot take advantage of due to my truncations, but the general shape of the distribution is suitable for Gaussian variances anyway. The “hyper-hyper-priors” α and β of the inverse-Gamma are chosen somewhat arbitrarily, but this is not a huge issue since bias in hyper-hyper-parameters have little influence over the posterior. I will discuss my decision strategy for these hyper-hyper-priors in a later subsection.

My final assumption is that the prior parameters are independent. This is actually a poor assumption that I will revisit in the discussion section, but I have no way to quantify the covariances of the parameters *a priori*, so I have to make this assumption to make modeling possible. Additionally, I make the assumption that the hyperparameters are independent from each other. This allows me to make some simplifications using conditional independence in my posterior formulation.

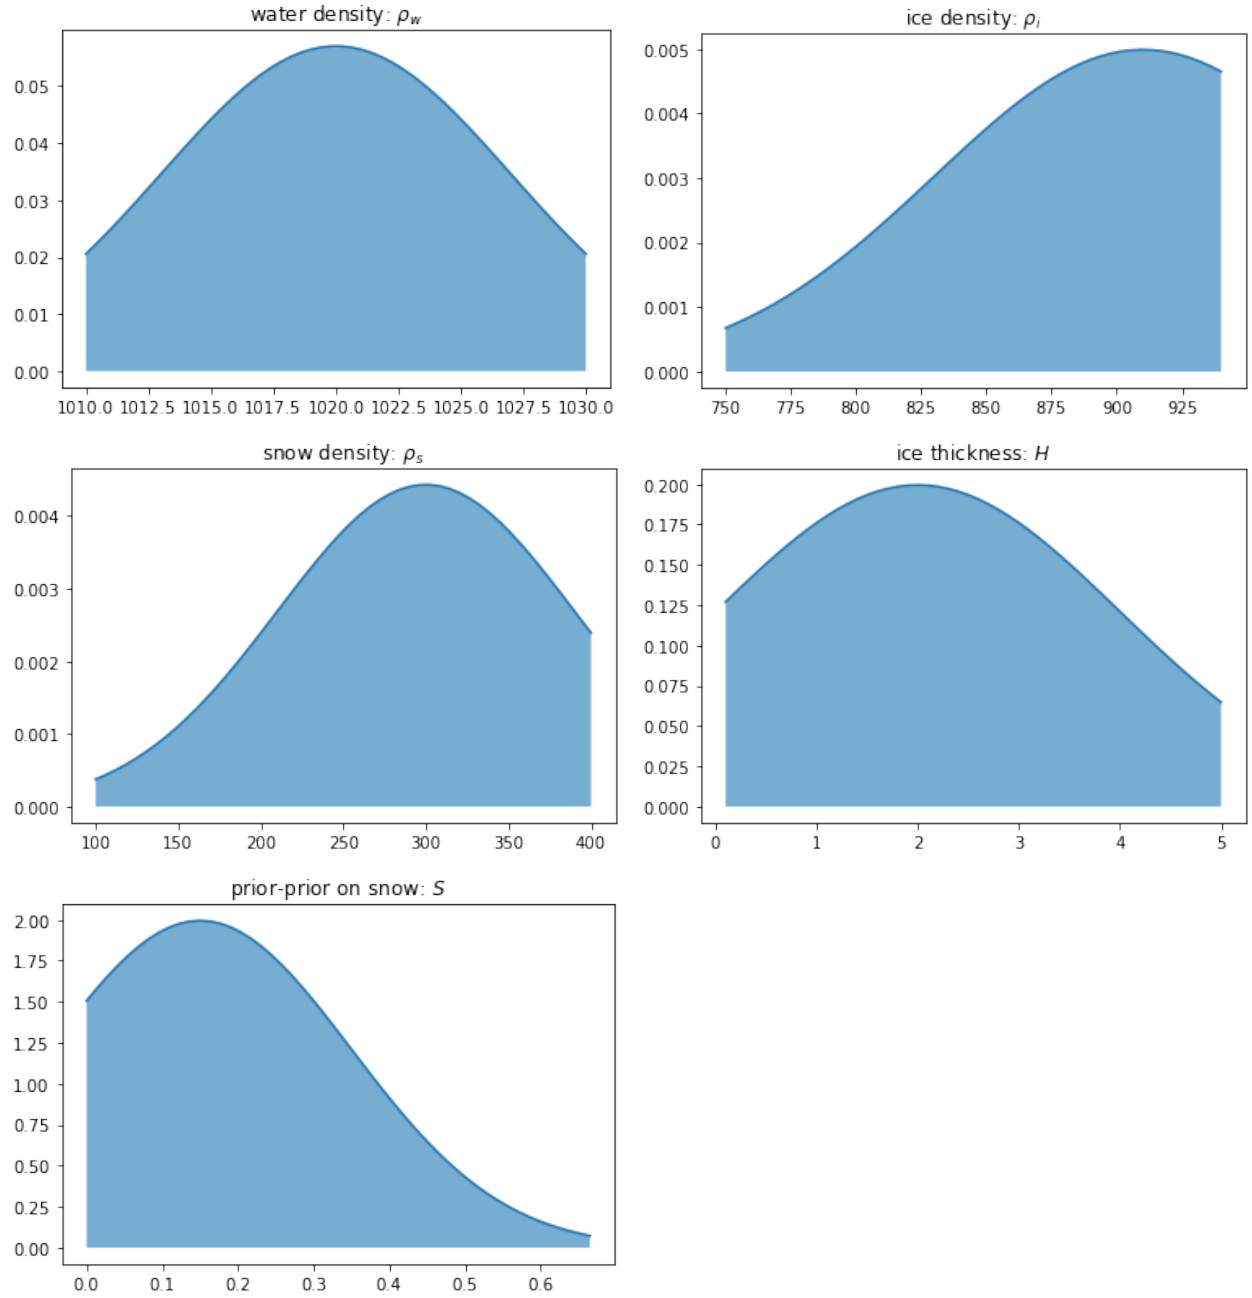
2.3 Prior Formulation and Densities

As I mentioned in the assumptions section, I chose the variances on the three densities rather arbitrarily. I estimated the standard deviations of ρ_w to 7, ρ_i to 80, and ρ_s to 90, and squared those to obtain the variances. I also set the standard deviation of ice density to 2 meters, equal to the mean, to account for a fair amount of uncertainty, and squared that as well. Finally, I guessed a mean standard deviation of snow density equal 0.2 meters, which is very wide. However, this value will not be going straight into the prior model, since I will be modeling that standard deviation with an inverse-Gamma hyperprior. It remains useful for getting situated and visualizing the prior, though.

parameter	mean	std. dev	variance	lower bound	upper bound
ρ_w	1020	7	49	1010	1030
ρ_i	910	80	6400	750	940
ρ_s	300	90	8100	100	400
H	2.0	2.0	4.0	0.1	5.0
S	0.15	≈ 0.2	≈ 0.04	0	unknown

Plots of the prior are shown on the next page. Note that the snow thickness plot is not a “true” prior; it is only an estimate based on the hyperprior I will define next.

Figure 1: Plots of Prior Densities



2.4 Hyperprior Formulation and Densities

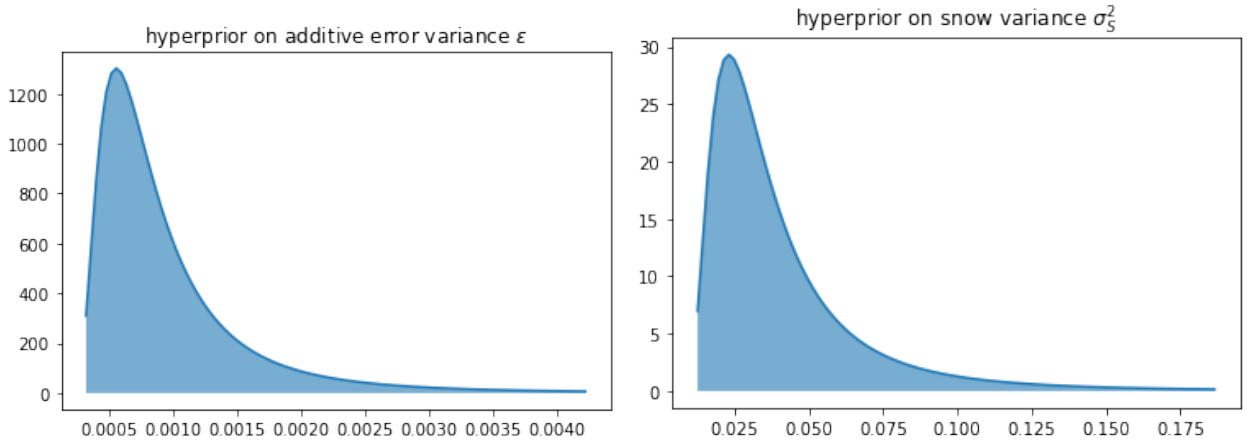
As mentioned in the assumptions, I chose to use inverse-Gamma densities as hyperpriors on the additive noise variance as well as the snow depth variance. The Gamma distribution contains two parameters in the typical parametrization: α and β . These are largely arbitrary for my usage, so I did a re-parametrization in terms of mean and variance, and used identities to convert them back to α and β (from Wikipedia):

$$\mu = \frac{\beta}{\alpha - 1} \text{ for } \alpha > 1; \sigma^2 = \frac{\beta^2}{(\alpha - 1)^2(\alpha - 2)} \text{ for } \alpha > 2$$

I chose hyperprior means equal to my estimates of the variances of the snow depth and error noise, which are the squared standard deviations. For snow depth, I chose a semi-arbitrary estimate of 0.2 meters as mentioned earlier. For additive noise, I estimated three centimeters based on the description of “accurate to a few centimeters” on IceSat-2. I chose the hyperprior variances assuming that the standard deviations would be equal to the means. After that, I used the sympy package in Python to solve the equations for the hyper-hyperparameters α and β . Additionally, I also added a small “nugget” to each distribution to shift them a tiny bit away from zero to aid computation.

hyperparameter	expected mean	std. dev	variance	α	β	nugget
σ_ϵ^2	$0.03^2 = 9.0\text{e-}4$	9e-4	8.1e-7	3.0	0.0018	0.0001
σ_S^2	$0.2^2 = 0.04$	0.04	1.6e-3	3.0	0.08	0.003

Figure 2: Plots of Hyperprior Densities



2.5 Forward Model and Likelihood Function

Since I assume an additive Gaussian noise error, the likelihood function is a simple extension of the physical forward model. Let x be the five-dimensional vector containing the parameters ρ_w , ρ_i , ρ_s , H , and S , and let $g(x)$ be the forward model predicting some freeboard value F_{pred} from parameters x . We can then describe the likelihood function as a Gaussian distribution with mean zero and variance σ_ϵ^2 , evaluated at the difference $F - g(x)$ as observed minus predicted.

$$f(F|x, \sigma_\epsilon^2) \sim N(0, \sigma_\epsilon^2)$$

This is notably dependent on a variable noise variance, which I will explore in the next subsection on the hierarchical model.

2.6 Hierarchical Model and Posterior Formulation

Now I will proceed to discuss the posterior of the model, which contains hierarchical components regarding the noise and snow variances. Traditionally, a non-hierarchical Bayesian model follows Bayes' theorem, the posterior being proportional to the prior times the likelihood:

$$f(x|F) \propto f(F|x)f(x)$$

However, we have a hierarchical model which tries to incorporate the observation and obtain posterior hyperparameters as well. These posterior hyperparameters show up here:

$$f(x, \sigma_\epsilon^2, \sigma_S^2|F) \propto f(F|x, \sigma_\epsilon^2)f(\sigma_\epsilon^2|\alpha_\epsilon, \beta_\epsilon)f(x|\sigma_S^2)f(\sigma_S^2|\alpha_S, \beta_S)$$

I have dropped a few terms in this equation due to conditional independence. For example, σ_S does not show up in the likelihood function because F is conditionally independent of σ_S given some value of x that has already been evaluated. I also make use of my assumption that the two hyperparameter variances are independent as well, to make the joint distribution of σ_ϵ^2 and σ_S^2 just equal to the one we are interested in when the other is held constant. Finally, I have left out the α and β terms on the left side to save space; they are just constants here and do not need to be listed, but I'm listing them to make clear which terms are inverse-Gamma hyperpriors.

I will elaborate on the specific computation of this posterior, as well as the hierarchical separation of the parameters and hyperparameters, in the next section.

2.7 Expected Value of Ice Thickness

The solution to my answer is relatively simple after computing the posterior density. Since I am purely making an inference on the parameters, I can just compute $\mathbb{E}[x_{post}]$ and then extract the specific expected value and posterior density on ice thickness H .

3 Solution and Computation Strategy

3.1 Non-Conjugacy and Markov Chain Monte Carlo

My model is non-conjugate in two areas. First, the physical model is simple but non-linear, so no linear algebra identities can be used to analytically derive the distributions. Second, I chose to use truncated Gaussians, which have complex normalizing constants that render conjugacy impossible in and of itself. These conditions leave sampling as the best technique to derive the posterior, as we can take advantage of the Central Limit Theorem to prevent the exponential growth in computational complexity for numerical integration of a multi-dimensional density.

Within sampling techniques, I chose to use Markov Chain Monte Carlo due to the difficulty of generating direct samples from the truncated prior density.

3.2 Hierarchical MCMC, Metropolis-Within-Gibbs Sampling

With a hierarchical model, it is easier to break up each sampling step into three distinct steps: one for the parameters x and one for each hyperparameter σ_ϵ^2 and σ_S^2 . This technique, called Gibbs sampling, is computationally easier to evaluate. In each sub-step, I can hold the other variables constant, reducing the amount of required computation. In addition, I free up the hyperparameters to be sampled independently of the parameters. Gibbs sampling will necessitate a modification to the formulation where each parameter and hyperparameter is isolated in the posterior.

For the sub-step involving the parameters x while the hyperparameters are held constant, the marginal posterior on x looks like this, which can then be simplified by dropping σ_S^2 in the likelihood function and σ_ϵ^2 in the prior due to our independence and conditional independence assumptions on the parameters.

$$\begin{aligned} f(x|F, \sigma_\epsilon^2, \sigma_S^2) &\propto f(F|x, \sigma_\epsilon^2, \sigma_S^2)f(x|\sigma_\epsilon^2, \sigma_S^2) \\ &\propto f(F|x, \sigma_\epsilon^2)f(x|\sigma_S^2) \end{aligned}$$

For the sub-step involving σ_ϵ^2 while x and σ_S^2 are held constant, the marginal posterior looks like this, which can be simplified by dropping dependence on σ_S^2 entirely (since I assume the two hyperpriors are independent), and dropping both dependencies on x and σ_S^2 in the hyperprior (again due to the same assumption).

$$\begin{aligned} f(\sigma_\epsilon^2|F, \sigma_S^2) &\propto f(F|x, \sigma_\epsilon^2, \sigma_S^2)f(\sigma_\epsilon^2|x, \sigma_S^2) \\ f(\sigma_\epsilon^2|F) &\propto f(F|x, \sigma_\epsilon^2)f(\sigma_\epsilon^2) \end{aligned}$$

Finally, for the sub-step involving σ_S^2 while x and σ_ϵ^2 are held constant, the marginal posterior looks like this:

$$\begin{aligned} f(\sigma_S^2|F, \sigma_\epsilon^2) &\propto f(F|x, \sigma_\epsilon^2, \sigma_S^2)f(\sigma_S^2|x, \sigma_\epsilon^2) \\ &\propto f(F|x, \sigma_\epsilon^2)f(x|\sigma_S^2)f(\sigma_S^2) \end{aligned}$$

In the σ_S^2 step, I first simplify by dropping σ_S^2 in the likelihood due to conditional independence given fixed x ,

3.3 Overview of Software Tools

4 Results

4.1 Sampling Results and Posterior Densities

4.2 Central Limit Theorem, Autocorrelation, and Error Analysis

5 Discussion

5.1 Sea Ice Thickness

5.2 Observations on Posterior

5.3 Revisiting Assumptions and Potential Bias